

Data Mining

Practical Machine Learning Tools and Techniques

Slides for Chapter 9, Probabilistic methods

of *Data Mining* by I. H. Witten, E. Frank,
M. A. Hall, and C. J. Pal

Random variables

- In probabilistic approaches to machine learning it is common to think of data as observations arising from an underlying probability model for *random variables*
- Given a discrete random variable A , $P(A)$ is a function that encodes the probabilities for each of the categories, classes or states that A may be in
- For a continuous random variable x , $p(x)$ is a function that assigns a probability density to all possible values of x
- In contrast, $P(A=a)$ is the single probability of observing the specific event $A=a$

Notation

- The $P(A=a)$ notation is often simplified to simply $P(a)$, but one must remember if a was defined as a random variable or as an observation
- Similarly for the observation that continuous random variable x has the value x_1 it is common to write this as $p(x_1)=p(x=x_1)$, a simplification of the longer but clearer notation

The product rule

- The *product rule*, sometimes referred to as the “fundamental rule of probability,” states that the joint probability of random variables A and B can be written

$$P(A, B) = P(A | B)P(B)$$

- The product rule also applies when A and B are groups or subsets of events or random variables.

The sum rule

- The *sum rule* states that given the joint probability of variables X_1, X_2, \dots, X_N , the *marginal probability* for a given variable can be obtained by summing (or integrating) over all the other variables.
- For example, to obtain the marginal probability of X_1 , sum over all the states of all the other variables:

$$P(X_1) = \sum_{x_2} \dots \sum_{x_N} P(X_1, X_2 = x_2, \dots, X_N = x_N)$$

Marginalization

- The previous notation can be simplified to

$$p(x_1) = \sum_{x_2} \dots \sum_{x_N} P(x_1, x_2, \dots, x_N)$$

- The sum rule generalizes to continuous random variables, ex. for x_1, x_2, \dots, x_N we have

$$p(x_1) = \dots \int_{x_N} p(x_1, x_2, \dots, x_N) dx_1 \dots dx_N$$

- These procedures are known as **marginalization**
- They give us **marginal distributions** of the variables not included in the sums or integrals

Bayes' Rule

- Can be obtained by swapping A and B in the product rule and observing $P(B|A)P(A)=P(A|B)P(B)$ and therefore

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

- Suppose we have models for $P(A|B)$ and $P(B)$
 - We observe that $A=a$, and
 - we want to compute $P(B|A=a)$
 - $P(A=a|B)$ is referred to as the **likelihood**
 - $P(B)$ is the *prior* distribution of B
 - $P(B|A=a)$ is posterior distribution, obtained from:

$$P(A = a) = \sum_b P(A = a, B = b) = \sum_b P(A = a | B = b)P(B = b)$$

Maximum Likelihood

- Our goal is to estimate a set of parameters θ of a probabilistic model, given a set of *observations* x_1, x_2, \dots, x_n .
- Maximum likelihood techniques assume that:
 - 1) the examples have no dependence on one another, the occurrence of one has no effect on the others, and
 - 2) each can be modeled in exactly the same way.
- These assumptions are often summarized by saying that events are *independent and identically distributed* (i.i.d.).

Maximum Likelihood

- The i.i.d. assumption corresponds to the use of a joint probability density function for all observations consisting of the product of the same probability model $p(x_i; \theta)$ applied to each observation independently.
- For n observations, this could be written as

$$p(x_1, x_2, \dots, x_n; \theta) = p(x_1; \theta)p(x_2; \theta) \dots p(x_n; \theta)$$

where each function $p(x_i; \theta)$ has the same θ

Maximum Likelihood

- The likelihood of our data can be written

$$L(\theta; x_1, x_2, \dots, x_n) = \prod_{i=1}^n p(x_i; \theta)$$

- The data is fixed, but we can adjust θ so as to *maximize the likelihood or log-likelihood*

$$\theta_{ML} = \arg \max_{\theta} \sum_{i=1}^n \log p(x_i; \theta)$$

- We use the the log-likelihood as it is more numerically stable

Maximum a posteriori (MAP) parameter estimation

- If we treat our parameters as random variables we can compute the posterior

$$p(\theta | x_1, x_2, \dots, x_n) = \frac{p(x_1, x_2, \dots, x_n | \theta) p(\theta)}{p(x_1, x_2, \dots, x_n)}$$

- We have used $|$ or the “given” notation in place of $;$ to emphasize that θ is random, but
- Conditioned on a point estimate for the posterior we have a conditionally i.i.d. model
- MAP parameter estimation seeks

$$\theta_{MAP} = \arg \max_{\theta} \left[\sum_{i=1}^n \log p(x_i; \theta) + p(\theta) \right]$$

The chain rule of probability

- Results from applying the product rule recursively between a single variable and the rest of the variables
- The **chain rule** states that the joint probability of n attributes $A_{i=1...m}$ can be decomposed into the following product:

$$P(A_1, A_2, \dots, A_n) = P(A_1) \prod_{i=1}^{n-1} P(A_{i+1} | A_i, A_{i-1}, \dots, A_1)$$

e.g.

$$P(A_1, A_2, A_3, A_4) = P(A_1) P(A_2 | A_1) P(A_3 | A_2, A_1) P(A_4 | A_3, A_2, A_1)$$

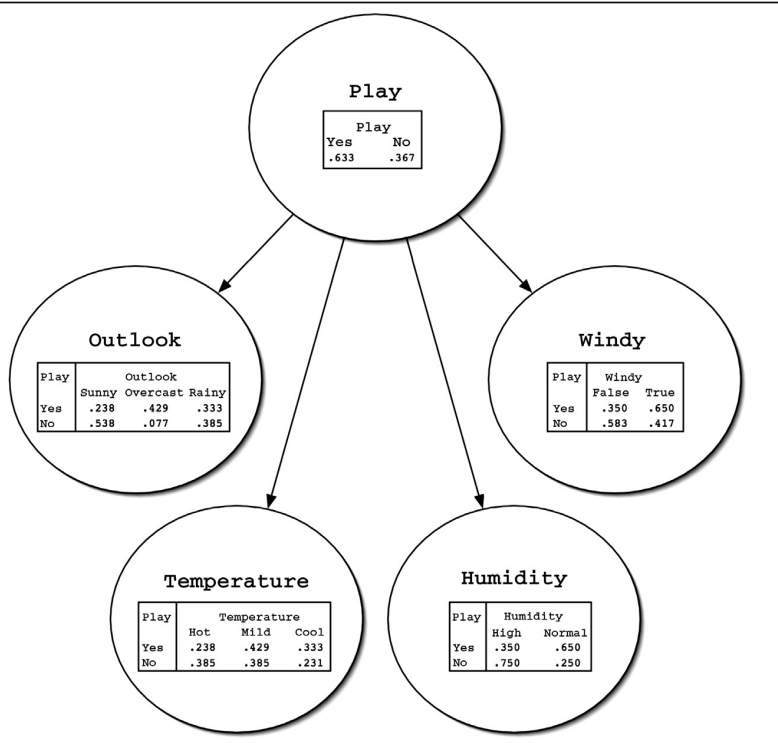
Bayesian networks

- The chain rule holds for any order for the A_i s
- A Bayesian network is an **acyclic** graph,
- Therefore its nodes can be given an **ordering** where ancestors of node A_i have indices $< i$
- Thus a Bayesian network can be written

$$P(A_1, A_2, \dots, A_n) = \prod_{i=1}^n P(A_i | \text{Parents}(A_i))$$

- When a variable has no parents, we use the unconditional probability of that variable

Bayesian network #1 for the weather data



Random Variables


Y: Play

O: Outlook

T: Temperature

H: Humidity

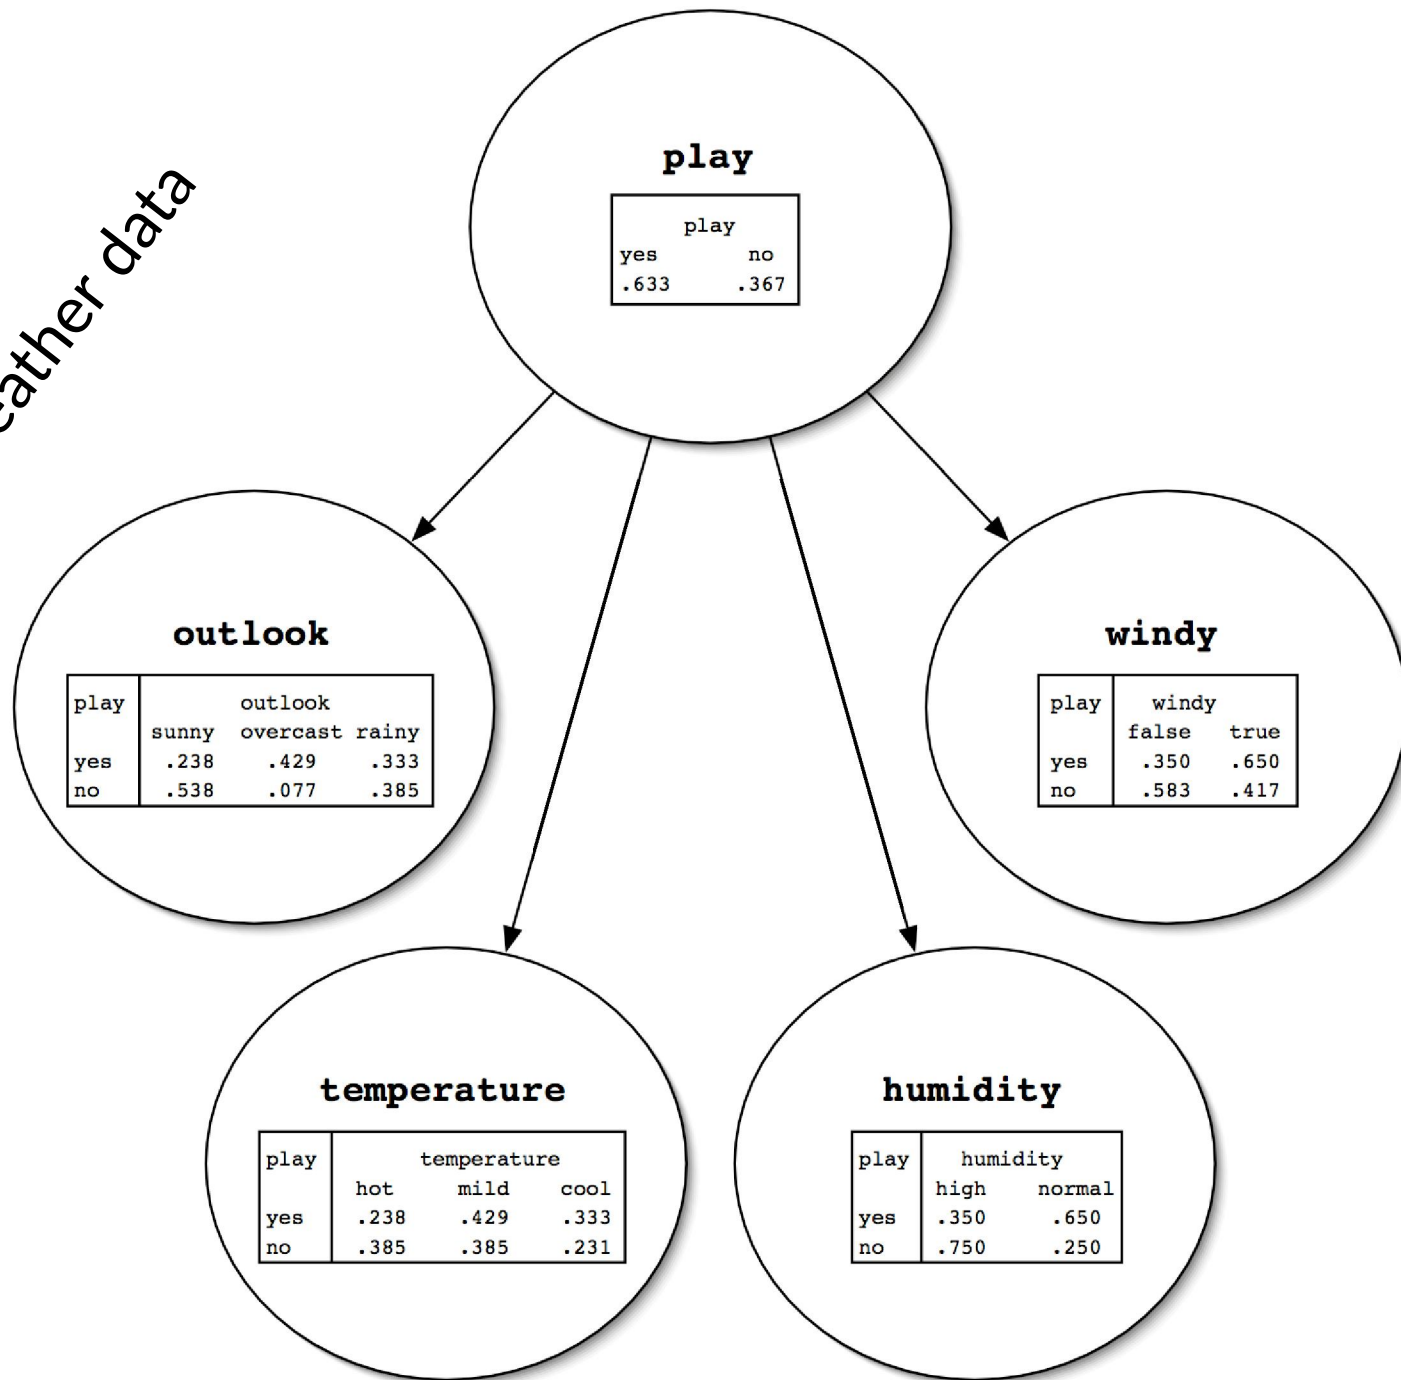
W: Windy



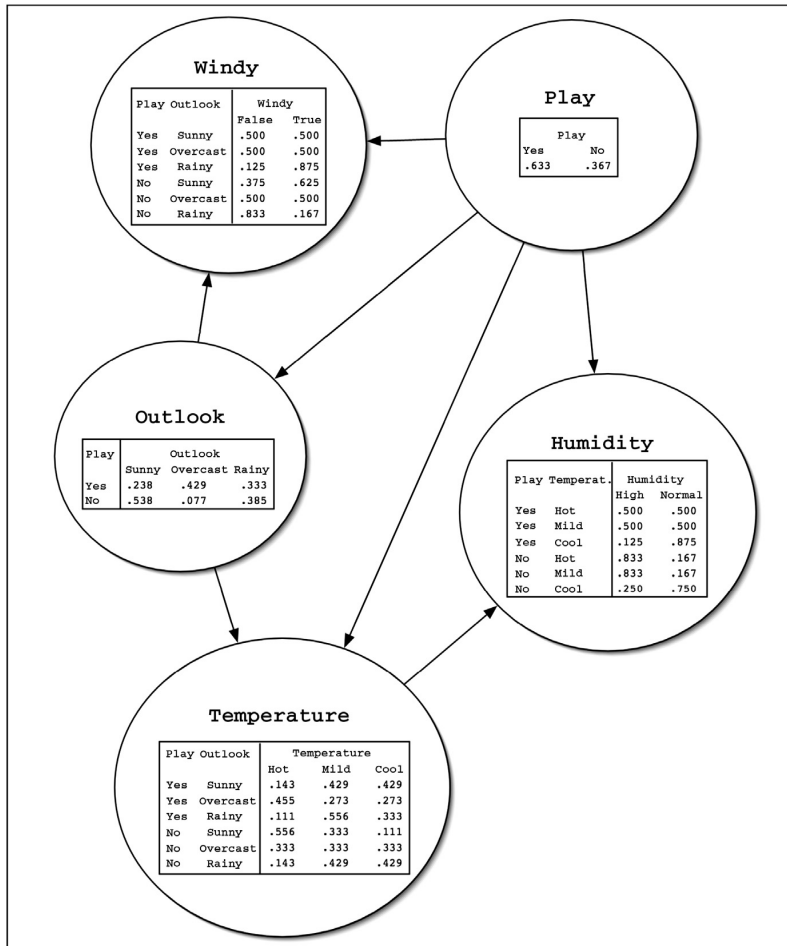
The graphs express the factorization below:

$$P(Y, O, T, H, W) = P(W | Y)P(O | Y)P(T | Y)P(H | Y)P(Y)$$

Network #1 for the weather data



Bayesian network #2 for the weather data



Random Variables


Y: Play

O: Outlook

T: Temperature

H: Humidity

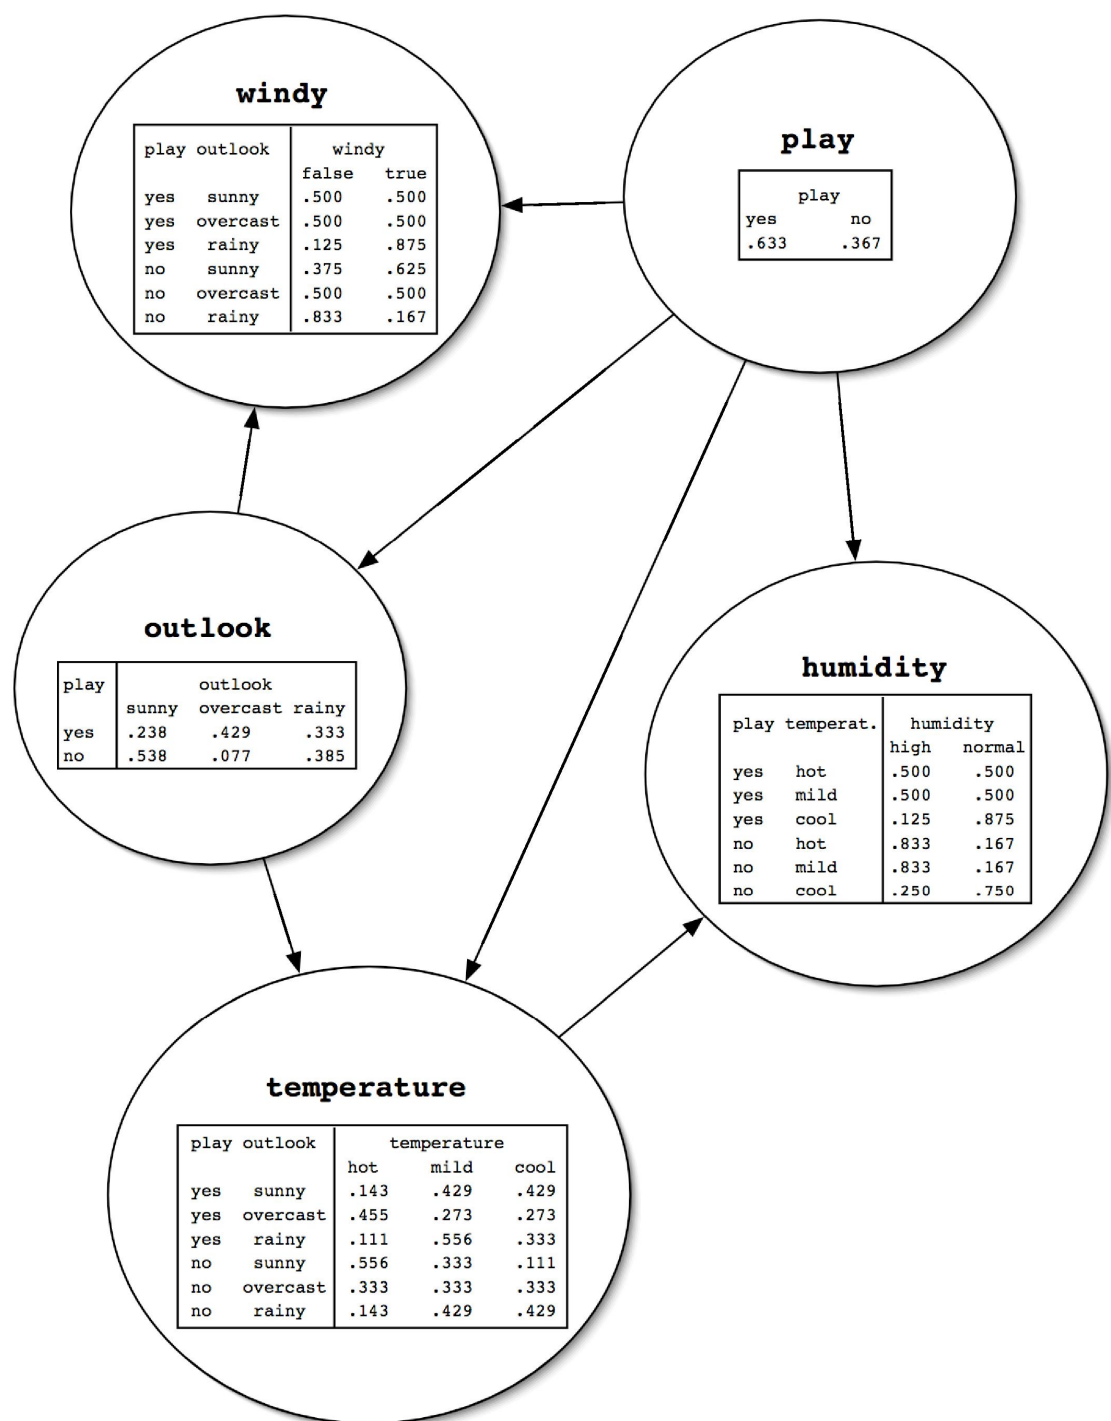
W: Windy



The graphs express the factorization below:

$$P(Y, O, T, H, W) = P(W | O, Y) P(O | Y) P(T | O, Y) P(H | T, Y) P(Y)$$

Network #2 for the weather data



Computing the class probabilities

- Two steps: computing a product of probabilities for each class and normalization
 - 1) For each class value
 - Take all attribute values and class value
 - Look up corresponding entries in conditional probability distribution tables
 - Take the product of all probabilities
 - 2) Divide the product for each class by the sum of the products (normalization)

Naive Bayes vs. Bayesian Network

Naive Bayes:

$P(\text{windy}=\text{true}, \text{outlook}=\text{sunny}, \text{temperature}=\text{cool}, \text{humidity}=\text{high} \mid \text{play}=\text{yes}) =$

$P(\text{windy}=\text{true} \mid \text{play}=\text{yes}) * P(\text{outlook}=\text{sunny} \mid \text{play}=\text{yes}) * \\ P(\text{temperature}=\text{cool} \mid \text{play}=\text{yes}) * P(\text{humidity}=\text{high} \mid \text{play}=\text{yes})$

Bayesian network:

$P(\text{windy}=\text{true}, \text{outlook}=\text{sunny}, \text{temperature}=\text{cool}, \text{humidity}=\text{high} \mid \text{play}=\text{yes}) =$

$P(\text{windy}=\text{true} \mid \text{outlook}=\text{sunny}, \text{play}=\text{yes}) * \\ P(\text{outlook}=\text{sunny} \mid \text{play}=\text{yes}) * \\ P(\text{temperature}=\text{cool} \mid \text{outlook}=\text{sunny}, \text{play}=\text{yes}) * \\ P(\text{humidity}=\text{high} \mid \text{temperature}=\text{cool}, \text{play}=\text{yes})$

Why can we do this? (Part I)

- Single assumption: values of a node's parents completely determine probability distribution for current node

$$Pr[\text{node}|\text{ancestors}] = Pr[\text{node}|\text{parents}]$$

- Means that node/attribute is conditionally independent of other ancestors given parents

Why can we do this? (Part II)

- Chain rule from probability theory:

$$Pr[a_1, a_2, \dots, a_n] = \prod_{i=1}^n Pr[a_i | a_{i-1}, \dots, a_1]$$

- Because of our assumption from the previous slide:

$$Pr[a_1, a_2, \dots, a_n] = \prod_{i=1}^n Pr[a_i | a_{i-1}, \dots, a_1] = \\ \prod_{i=1}^n Pr[a_i | a_i \text{'s parents}]$$

Chain Rule vs. Bayesian Network

- **Chain rule:**

$$P(a_1, a_2, a_3, a_4) = P(a_1 | a_2, a_3, a_4) * P(a_2, a_3, a_4)$$


$$P(a_2, a_3, a_4) = P(a_2 | a_3, a_4) * P(a_3, a_4)$$

$$P(a_3, a_4) = P(a_3 | a_4) * P(a_4)$$

$$P(a_1, a_2, a_3, a_4) = P(a_1 | a_2, a_3, a_4) * P(a_2 | a_3, a_4) * \\ P(a_3 | a_4) * P(a_4)$$

- **Bayesian network**

$$P(a_1, a_2, a_3, a_4) = \\ P(a_1 | a_4) * P(a_2 | a_3, a_4) * P(a_3 | a_4) * P(a_4)$$



Learning Bayes nets

- Basic components of algorithms for learning Bayes nets:
 - 1) Method for evaluating the goodness of a given network
 - Measure based on probability of training data given the network (or the logarithm thereof)
 - 2) Method for searching through space of possible networks
 - Amounts to searching through sets of edges because nodes are fixed

Estimating Bayesian network parameters

- The log-likelihood of a Bayesian network with V variables and N examples of complete variable assignments to the network is

$$\sum_{i=1}^N \log P(\{\tilde{A}_1, \tilde{A}_2, \dots, \tilde{A}_V\}_i) = \sum_{i=1}^N \sum_{v=1}^V \log P(\tilde{A}_{v,i} \mid \text{Parents}(\tilde{A}_{v,i}); \Theta_v)$$

where the parameters of each conditional or unconditional distribution are given by Θ_v

- We use the $\tilde{A}_{v,i}$ notation to indicate the i th observation of variable v

Estimating probabilities in Bayesian networks

- The estimation problem *decouples* into separate estimation problems for each conditional or unconditional probability
- Unconditional probabilities can be written as

$$P(A = a) = \frac{1}{N} \sum_{i=1}^N \mathbf{1}(\tilde{A}_i = a)$$

where $\mathbf{1}(\tilde{A}_i = a)$ is an indicator function returning 1 when the i^{th} observed value for $A_i = a$ and 0 otherwise

Estimating conditional distributions

- Estimating conditional distributions in Bayesian networks is equally easy and amounts to simply counting configurations and dividing, ex.

$$P(B = b \mid A = a) = \frac{P(B = b, A = a)}{P(A = a)} = \frac{\sum_{i=1}^N \mathbf{1}(\tilde{A}_i = a, \tilde{B}_i = b)}{\sum_{i=1}^N \mathbf{1}(\tilde{A}_i = a)}.$$

- Zero counts cause problems and this motivates the use of Bayesian priors

Estimating network structure

- One possibility is to use **cross-validation** to estimate the goodness of fit on held out data (so as to avoid over fitting) another is to **penalize model complexity**
- Let K be the number of parameters, LL the log-likelihood, and N the number of instances in the data.
- Two popular measures for evaluating the quality of a network are the **Akaike Information Criterion (AIC)**:

$$\text{AIC score} = -LL + K$$


and the following **MDL** metric based on the MDL principle:

$$\text{MDL score} = -LL + \frac{K}{2} \log N$$

- In both cases the log-likelihood is negated, so the aim is to minimize these scores.
- More Bayesian approach: use prior over model structures


The K2 Network structure learning algorithm

- **K2**: a simple and very fast learning algorithm,
- Starts with a given ordering of the nodes
- Processes each node in turn and greedily considers adding edges from previously processed nodes to the current one
- In each step it adds the edge that maximizes the network's score
- When there is no further improvement, attention turns to the next node
- The number of parents for each node can be restricted to a predefined maximum to mitigate overfitting



Tree augmented naïve bayes (TAN)

- Another good learning algorithm for Bayesian network classifiers
- Takes the Naïve Bayes (NB) classifier and adds edges to it
- The class attribute is the sole parent of each node in a NB model: TAN considers adding a second parent to each node
- If the class node and all corresponding edges are excluded from consideration, and assuming that there is exactly one node to which a second parent is not added, the resulting classifier has a tree structure rooted at the parentless node—hence the name
- For this restricted network type there is an efficient algorithm based on computing a maximum weighted spanning tree
- Algorithm is linear in the number of instances and quadratic in the number of attributes

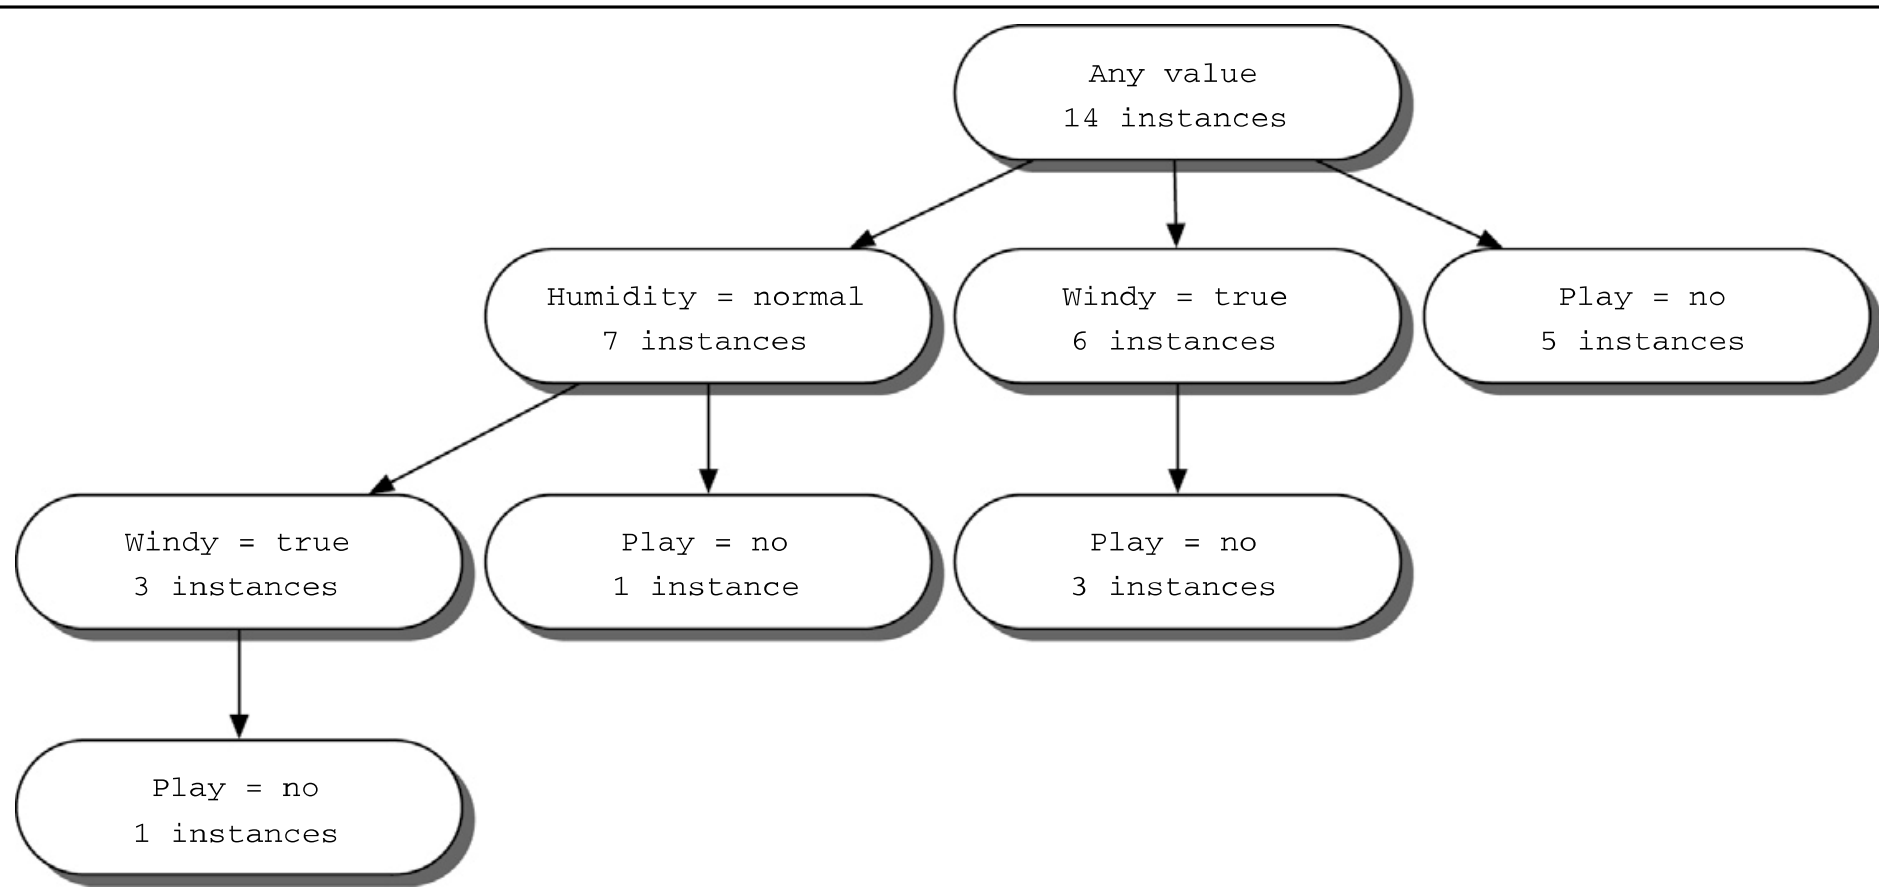


Data structures for fast Learning

- Learning Bayesian networks involves a lot of counting
- For each network structure considered in the search, the data must be scanned afresh to obtain the counts needed to fill out the conditional probability tables
- Counts can be stored effectively in a structure called an *all-dimensions (AD) tree*, which is analogous to the *kD*-trees used for nearest neighbor search
- Each node of the tree represents the occurrence of a particular combination of attribute values
- Straightforward to retrieve the count for a combination that occurs in the tree
- However, the tree does not explicitly represent many nonzero counts because the most populous expansion for each attribute is omitted

AD Tree example

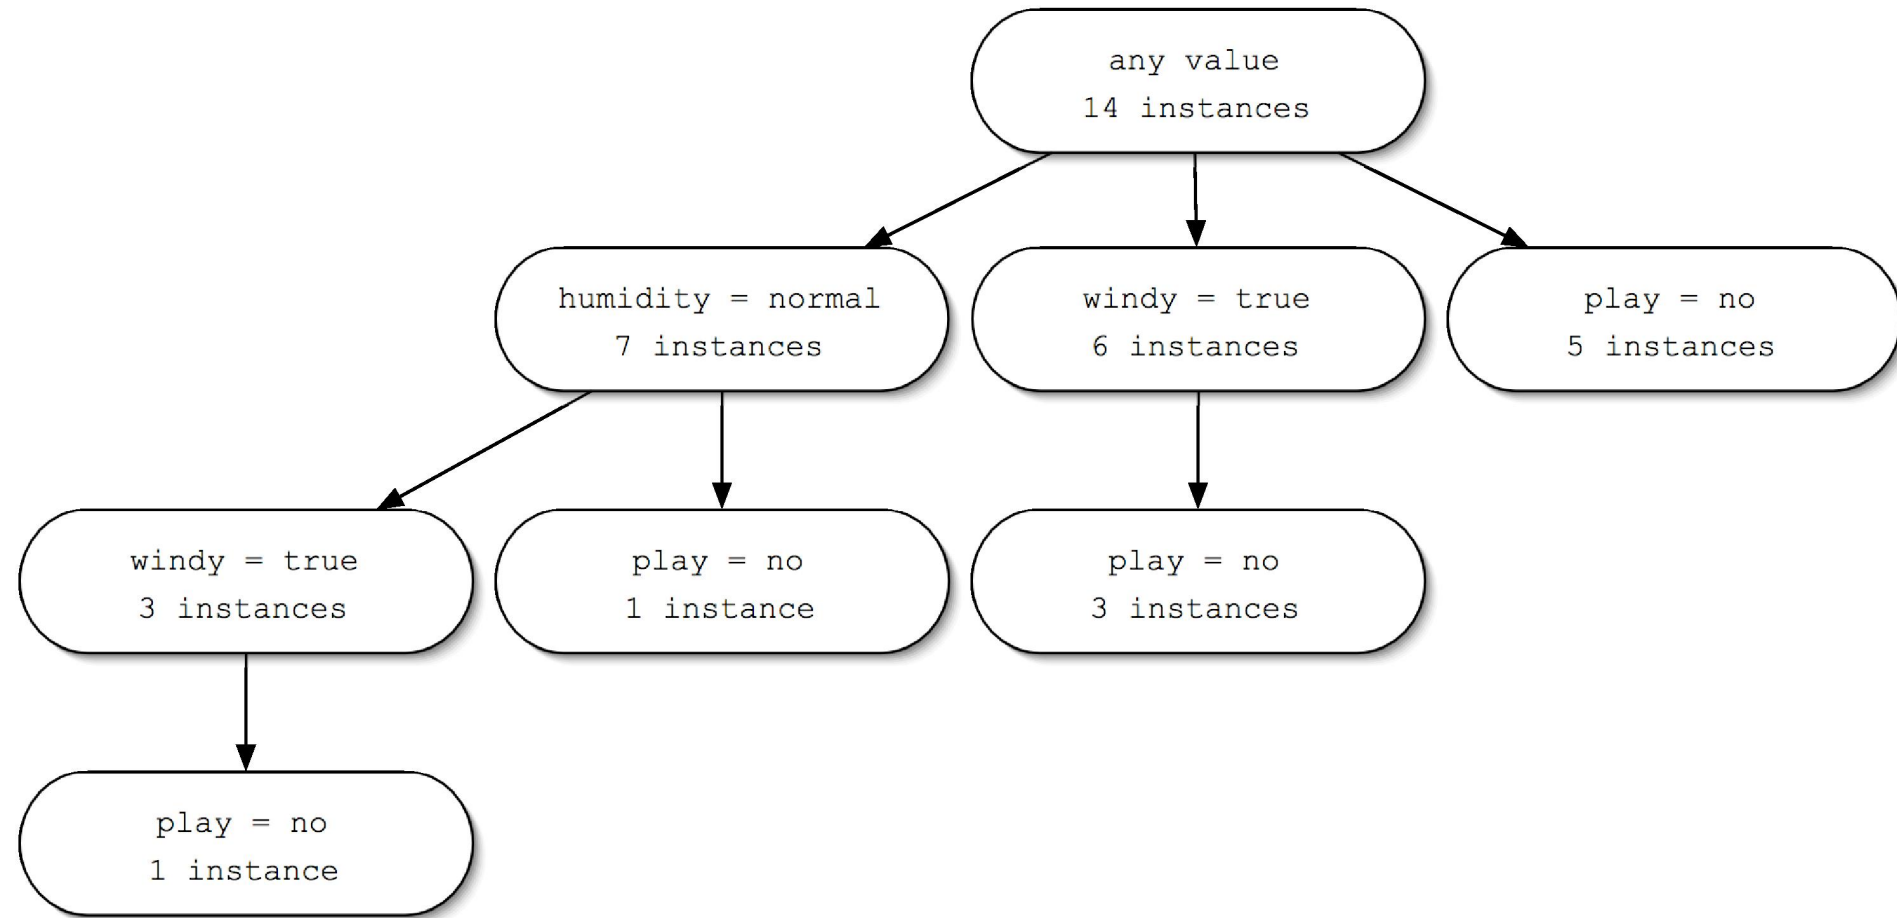
Humidity	Windy	Play	Count
High	True	Yes	1
High	True	No	2
High	False	Yes	2
High	False	No	2
Normal	True	Yes	2
Normal	True	No	1
Normal	False	Yes	4
Normal	False	No	0



Building an AD tree

- Assume each attribute in the data has been assigned an index
- Then, expand node for attribute i with the values of all attributes $j > i$
 - Two important restrictions:
 - Most populous expansion for each attribute is omitted (breaking ties arbitrarily)
 - Expansions with counts that are zero are also omitted
- The root node is given index zero

AD tree example



humidity = normal, windy = true, play = yes?

humidity = high, windy = true, play = no?

$\#(\text{windy}=\text{true}, \text{play}=\text{no} (3)) - \#(\text{humidity}=\text{normal}, \text{windy}=\text{true}, \text{play}=\text{no} (1))$ 33

Bibliographic Notes & Further Reading

Learning and Bayesian Networks

- The K2 algorithm for learning Bayesian networks was introduced by Cooper and Herskovits (1992).
- Bayesian scoring metrics are covered by Heckerman et al. (1995).
- Friedman et al. (1997) introduced the tree augmented Naïve Bayes algorithm, and also describe multinets.
- AD trees were introduced and analyzed by Moore and Lee (1998)
- Komarek and Moore (2000) introduce AD trees for incremental learning that are also more efficient for datasets with many attributes.


Clustering with a Gaussian Mixture

- Given the data on the left *without the labels A and B*, we wish to estimate a model for a two class **Gaussian Mixture Model (GMM)** on the right

Data

A 51	A 46	B 65
A 43	B 64	A 46
B 62	A 51	A 39
B 64	A 52	B 62
A 45	B 62	B 64
A 42	A 49	A 52
A 46	A 48	B 63
A 45	B 62	B 64
A 45	A 43	A 48
B 62	A 40	B 64
A 47	A 48	A 48
A 52	B 64	A 51
B 64	A 51	A 48
A 51	B 63	B 64
B 65	A 43	A 42
A 48	B 65	A 48
A 49	B 66	A 41

Model



Estimating Gaussian parameters

- If we knew which of the two distributions each instance came from, finding the five parameters would be easy—just estimate the mean and standard deviation for $n=n_A$ or $n=n_B$ samples x_1, x_2, \dots, x_n for each cluster, A and B

$$\mu = \frac{x_1 + x_2 + \dots + x_n}{n}$$

$$\sigma^2 = \frac{(x_1 - \mu)^2 + (x_2 - \mu)^2 + \dots + (x_n - \mu)^2}{n - 1}$$

- To estimate the fifth parameter p_A , just take the proportion of the instances that are in the A cluster, then $p_B = 1 - p_A$.

Motivating the EM algorithm

- If you knew the five parameters, finding the (posterior) probabilities that a given instance comes from each distribution would be easy
- Given an instance x_i , the probability that it belongs to cluster A is

$$P(A|x_i) = \frac{P(x_i | A) \cdot P(A)}{P(x_i)} = \frac{N(x_i; \mu_A, \sigma_A) p_A}{N(x_i; \mu_A, \sigma_A) p_A + N(x_i; \mu_B, \sigma_B) p_B}$$

where $N()$ is the normal or Gaussian distribution

$$N(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Learning the clusters

- Assume:
 - we know there are k clusters
- Learn the clusters \Rightarrow
 - determine their parameters
 - i. e. means and standard deviations
- Performance criterion:
 - *probability of training data given the clusters*
- EM algorithm
 - finds a local maximum of the likelihood

EM algorithm

- EM = Expectation-Maximization
 - Generalize k -means to probabilistic setting
- Iterative procedure:
 - E “expectation” step:
Calculate cluster probability for each instance
 - M “maximization” step:
Estimate distribution parameters from cluster probabilities
- Store cluster probabilities as instance weights
- Stop when improvement is negligible

More on EM

- Estimate parameters from weighted instances





$$\mu_A = \frac{w_1 x_1 + w_2 x_2 + \dots + w_n x_n}{w_1 + w_2 + \dots + w_n}$$

$$\sigma_A^2 = \frac{w_1 (x_1 - \mu)^2 + w_2 (x_2 - \mu)^2 + \dots + w_n (x_n - \mu)^2}{w_1 + w_2 + \dots + w_n}$$



- Stop when log-likelihood saturates
- Log-likelihood:

$$\sum_i \log(p_A \text{Pr}[x_i | A] + p_B \text{Pr}[x_i | B])$$

EM Example (1)



EM Example (2)



Extending the mixture Model

- The **Gaussian distribution** generalizes to n-dimensions
- Consider a two-dimensional model consisting of independent Gaussian distributions for each dimension
- We can transform from **scalar** to **matrix notation** for a two dimensional Gaussian distribution as follows:

$$\begin{aligned} P(x_1, x_2) &= \frac{1}{\sqrt{2\pi}\sigma_1} \exp\left[-\frac{(x_1 - \mu_1)^2}{2\sigma_1^2}\right] \frac{1}{\sqrt{2\pi}\sigma_2} \exp\left[-\frac{(x_2 - \mu_2)^2}{2\sigma_2^2}\right] \\ &= (2\pi)^{-1} (\sigma_1^2 \sigma_2^2)^{-1/2} \exp\left\{-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right\} \\ &= (2\pi)^{-1} |\boldsymbol{\Sigma}|^{-1/2} \exp\left\{-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right\}, \end{aligned}$$

- $\boldsymbol{\Sigma}$ is the covariance *matrix*, $|\boldsymbol{\Sigma}|$ is its determinant, the vector $\mathbf{x} = [x_1 \ x_2]^T$, and the mean vector $\boldsymbol{\mu} = [\mu_1, \mu_2]^T$

The multivariate Gaussian distribution

- Can be written in the following general form

$$P(x_1, x_2, \dots, x_d) = (2\pi)^{-d/2} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}.$$

- The equation for estimating the covariance matrix is

$$\Sigma = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \boldsymbol{\mu})(\mathbf{x}_i - \boldsymbol{\mu})^T.$$

- The mean is simply

$$\boldsymbol{\mu} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i.$$

Clustering with correlated attributes

- If all attributes are continuous one can simply use a **full covariance Gaussian mixture model**
- But one needs to estimate $n(n + 1)/2$ parameters per mixture component for a full covariance matrix model
- As we will see later principal component analysis (PCA) can be formulated as a probabilistic model, yielding probabilistic principal component analysis (PPCA),
- Approaches known as ***mixtures of principal component analyzers*** or ***mixtures of factor analyzers*** provide ways of using a much smaller number of parameters to represent large covariance matrices